

# Python and the AI Revolution

Chase Coleman and John Stachurski

March 2024

# Topics

We will discuss

- AI-driven scientific computing
- Where are we heading?
- Economic applications?

Sides

[https://github.com/QuantEcon/imf\\_2024](https://github.com/QuantEcon/imf_2024)

# AI-driven scientific computing

AI is changing the world

- deep learning / other machine learning
- large language models
- computer vision
- speech recognition
- scientific knowledge discovery
- forecasting and prediction, etc., etc.

The huge amount of resources being poured into AI is changing the choice set for **all** scientific coders

## Key players

- OpenAI / Microsoft
- Google (Google Research, Google DeepMind)
- Meta
- Anthropic, etc.

## Platforms / libraries

- PyTorch (ChatGPT, Meta's LLaMA 2, Stable Diffusion)
- Google JAX (Google's Gemini)
- Tensorflow, Keras, Mojo?

# Lightening introduction to deep learning

Supervised deep learning: find a good approximation to an unknown functional relationship

$$y = f(x)$$

- $x$  is the input and  $y$  is the output

## Examples.

- $x$  = unfinished sentence,  $y$  = next word
- $x$  = weather sensor data,  $y$  = max temp tomorrow

# Training

Nonlinear regression: Take data set  $(x_i, y_i)_{i=1}^n$  and solve

$$\min_{\theta} \ell(\theta) = \sum_{i=1}^n (y_i - f_{\theta}(x_i))^2 \quad \text{s.t.} \quad \theta \in \Theta$$


In the case of ANNs, we consider all  $f_{\theta}$  having the form

$$f_{\theta} = \sigma \circ A_1 \circ \cdots \circ \sigma \circ A_k$$


where

- $A_i x = W_i x + b_i$  is an affine map
- $\sigma$  is a nonlinear “activation” function

## Minimizing a smooth loss functions



Source: <https://danielkhv.com/>



Source: <https://losslandscape.com/gallery/>



## Core elements

- automatic differentiation!
- parallelization (CPUs / GPUs / TPUs)!
- Compilers / JIT-compilers!

---

```
import jax.numpy as jnp
from jax import grad, jit
```

```
def predict(params, x):
    for W, b in params:
        y = jnp.dot(W, x) + b
        x = jnp.tanh(y)
    return y
```

```
def loss(params, x, targets):
    preds = predict(params, x)
    return jnp.sum((preds - targets)**2)
```

```
grad_loss = jit(grad(loss))
```

*# Now use gradient descent on the loss function*

---

“ECMWF’s weather forecasting model is considered the gold standard for medium-term weather forecasting...Google DeepMind claims in an non-peer-reviewed paper to have beat it 90% of the time...”

“Traditional forecasting models are big, complex computer algorithms based on atmospheric physics and take hours to run. AI models can create forecasts in just seconds.”

Source: MIT Technology Review

## Relevant to economics?

Deep learning provides massively powerful pattern recognition

But macroeconomic data is

- far more limited than weather observation sensor data
- generally nonstationary
- laws of motion change with policies (Lucas critique)

Possible applications:

- Finding stylized facts? Testing causal relationships?
- Numerical methods – approximating high-dimensional functions

Relevant to economics?

Deep learning provides massively powerful pattern recognition

But macroeconomic data is

- far more limited than weather observation sensor data
- generally nonstationary
- laws of motion change with policies (Lucas critique)

Possible applications:

- Finding stylized facts? Testing causal relationships?
- Numerical methods – approximating high-dimensional functions

## One point of view

- Deep learning is not very relevant for policy-centric macroeconomic modeling
- Deep learning is yet to prove itself as a “better” approach to numerical methods
- And yet, at the same time, **the AI computing revolution is generating tools that are enormously beneficial for macroeconomic modeling**
  - autodiff, JIT compilers, parallelization, GPUs, etc.
- We can take full advantage of them right now...